

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
titanic_data=pd.read_csv('/content/titanic_train.csv')
len(titanic_data)
titanic_data.head()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2834
				Heikinen						

Next steps:

[Generate code with titanic\\_data](#)

[View recommended plots](#)

```
titanic_data.index
titanic_data.columns

Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
      'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
      dtype='object')
```


```
titanic_data.info()
titanic_data.describe()
```

<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 891 entries, 0 to 890  
Data columns (total 12 columns):  
# Column Non-Null Count Dtype  
--- ---  
0 PassengerId 891 non-null int64  
1 Survived 891 non-null int64  
2 Pclass 891 non-null int64  
3 Name 891 non-null object  
4 Sex 891 non-null object  
5 Age 714 non-null float64  
6 SibSp 891 non-null int64  
7 Parch 891 non-null int64  
8 Ticket 891 non-null object  
9 Fare 891 non-null float64  
10 Cabin 204 non-null object  
11 Embarked 889 non-null object  
dtypes: float64(2), int64(5), object(5)  
memory usage: 83.7+ KB

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```
sns.countplot(x='Survived',data=titanic_data)
```

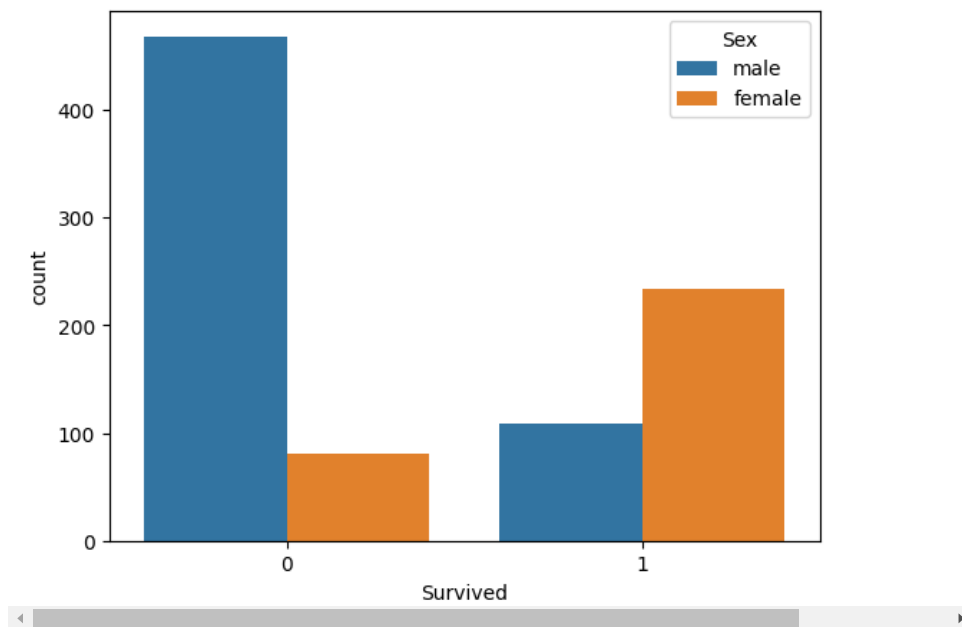
```
<Axes: xlabel='Survived', ylabel='count'>
```



```
sns.countplot(x='Survived',data=titanic_data,hue='Sex')  
titanic_data.isna()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
0		False	False	False	False	False	False	False	False	False	True
1		False	False	False	False	False	False	False	False	False	False
2		False	False	False	False	False	False	False	False	False	True
3		False	False	False	False	False	False	False	False	False	False
4		False	False	False	False	False	False	False	False	False	True
...	...	...	...	...	...	...	...	...	...	...	...
886		False	False	False	False	False	False	False	False	False	True
887		False	False	False	False	False	False	False	False	False	False
888		False	False	False	False	True	False	False	False	False	True
889		False	False	False	False	False	False	False	False	False	False
890		False	False	False	False	False	False	False	False	False	True

891 rows × 12 columns




```
titanic_data.isna().sum()
```

```

PassengerId    0
Survived        0
Pclass         0
Name           0
Sex            0
Age           177
SibSp          0
Parch          0
Ticket         0
Fare           0
Cabin         687
Embarked       2
dtype: int64


```

```
sns.heatmap(titanic_data.isna())
```



```
(titanic_data['Age'].isna().sum()/len(titanic_data['Age']))*100
(titanic_data['Cabin'].isna().sum()/len(titanic_data['Cabin']))*100
sns.displot(x='Age',data=titanic_data)
```


<seaborn.axisgrid.FacetGrid at 0x7ee58515d810>



```
titanic_data['Age'].fillna(titanic_data['Age'].mean(),inplace=True)
titanic_data['Age'].isna().sum()
```

0

```
sns.heatmap(titanic_data.isna())
```



```
titanic_data.drop('Cabin',axis=1,inplace=True)
titanic_data.head()
titanic_data.info()
titanic_data.dtypes
gender=pd.get_dummies(titanic_data['Sex'],drop_first=True)
titanic_data['Gender']=gender
titanic_data.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age          891 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(4)
memory usage: 76.7+ KB
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0		1	0	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1		2	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2834
				Heikkinen						

Next steps: [Generate code with titanic\\_data](#)

[View recommended plots](#)

```
titanic_data.drop(['Name', 'Sex', 'Ticket', 'Embarked'], axis=1, inplace=True)
titanic_data.head()
x=titanic_data[['PassengerId', 'Pclass', 'Age', 'SibSp', 'Parch', 'Fare', 'Gender']]
y=titanic_data['Survived']
y
```

```
0      0
1      1
2      1
3      1
4      0
..
886    0
887    1
888    0
889    1
890    0
```

```
from sklearn.metrics import confusion_matrix
pd.DataFrame(confusion_matrix(y_test, predict), columns=['Predicted No', 'Predicted Yes'], index=['Actual No', 'Actual Yes'])
```

	Predicted No	Predicted Yes	
Actual No	151	24	
Actual Yes	37	83	

```
#import train test split method
from sklearn.model_selection import train_test_split
#train test split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.33, random_state=42)
#import Logistic Regression
from sklearn.linear_model import LogisticRegression
#Fit Logistic Regression
lr=LogisticRegression()
lr.fit(x_train, y_train)
LogisticRegression()
#predict
predict=lr.predict(x_test)
from sklearn.metrics import classification_report
print(classification_report(y_test, predict))
from sklearn.metrics import confusion_matrix
pd.DataFrame(confusion_matrix(y_test, predict), columns=['Predicted No', 'Predicted Yes'], index=['Actual No', 'Actual Yes'])
```

	precision	recall	f1-score	support
0	0.80	0.86	0.83	175
1	0.78	0.69	0.73	120
accuracy			0.79	295
macro avg	0.79	0.78	0.78	295
weighted avg	0.79	0.79	0.79	295

/usr/local/lib/python3.10/dist-packages/sklearn/linear\_model/\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>